

# SC-2 Electric Boogalo

Regression on the Irish dataset aggregated by class

Kieran Morris, Cecina Babich Morrow and Sherman Kjo

## Contents

<b>1</b>	<b>Cleaning the Data</b>	<b>2</b>
1.1	Data overview . . . . .	2
1.2	Social class . . . . .	2
1.3	Loading and Structure . . . . .	2
1.4	Feature engineering . . . . .	2
1.4.1	Time-related features . . . . .	3
1.4.2	Temperature . . . . .	3
1.4.3	Fourier terms . . . . .	3
1.5	Scaling . . . . .	4
<b>2</b>	<b>Simple Regression</b>	<b>4</b>
2.1	Theory . . . . .	4
2.2	Model . . . . .	4
2.3	Implementation . . . . .	4
<b>3</b>	<b>Ridge Regression</b>	<b>4</b>
3.1	Theory . . . . .	4
3.2	Model . . . . .	4
3.3	Implementation . . . . .	4
<b>4</b>	<b>Gaussian Process Regression</b>	<b>4</b>
4.1	Theory . . . . .	4
4.2	Model . . . . .	5
4.3	Implementation . . . . .	5
<b>5</b>	<b>Results</b>	<b>7</b>

# 1 Cleaning the Data

## 1.1 Data overview

We are analyzing a set of Irish household electricity demand available from the `electBook` package. We have three datasets:

- `indCons`: 16799 x 2672 matrix of individual household electricity consumption. Each column corresponds to a household and each row to a time point. Demand is observed every half hour, so there are 48 observations per day per household.
- `survey`: 2672 row dataframe of household survey data. This dataset contains household level data on variables such as social class, renting vs. owning, appliances, etc.
- `extra`: 16799 row dataframe of time-related variables. This dataset contains the date-time of each demand observation, time of year, day of week, time of day, whether the day was a holiday, and external temperature.

```
# Extract individual dataframes
library(electBook)
library(tidyverse)
data(Irish)
indCons <- Irish[["indCons"]]
survey <- Irish[["survey"]]
extra <- Irish[["extra"]]
```

## 1.2 Social class

We wanted to investigate demand patterns across different social classes. The dataset includes 5 social classes, defined by the occupation of the head of household:

- AB: managerial roles, administrative or professional
- C1: supervisory, clerical, junior managerial
- C2: skilled manual workers
- DE: semi-skilled and unskilled manual workers, state pensioners, casual workers
- F: farmers

**\*\* Insert plot of demand patterns for different classes \*\***

We modeled the average demand for each social class separately.

## 1.3 Loading and Structure

We cleaned the data using the code available in `data_cleaning.R`, which writes out three dataframes: `df_halfhr`, `df_hr`, and `df_day`. These dataframes contain the half-hourly, hourly, and daily average demand data for each social class, as well as temperature (also aggregated to the relevant time scale), time of year (the time of year with 1st January represented as 0 and 31st December represented as 1), and day of the week. These dataframes are written to the `data` folder.

## 1.4 Feature engineering

Based on exploratory data analysis, we created some features from the dataset to model demand. Feature engineering (with the exception of the addition of the Fourier terms) was performed in the `feature_engineering.R` script, and the resulting dataframes `df_halfhr_scaled`, `df_hr_scaled`, and `df_day_scaled` were saved to the `data` folder.

### 1.4.1 Time-related features

We extracted the hour of the day and the month from the date-time variable. We also included a quadratic term for temperature to capture the non-linear relationship between temperature and demand. We also used one-hot encoding to include the day of the week in our models.

### 1.4.2 Temperature

We can visualize the relationship between temperature and the aggregate demand over time across all households:

```
agg <- colMeans(indCons)

temp_demand <- data.frame(demand = agg) %>%
  bind_cols(Irish[["extra"]])

ggplot(temp_demand, aes(x = dateTime, y = demand, color = temp)) +
  geom_point() +
  viridis::scale_color_viridis(option = "magma") +
  labs(x = "Date", y = "Total Demand", color = "Temperature") +
  theme_bw()
```

We can see that during the warmer summer months, demand dips, although the pattern is messy. We included linear and quadratic terms for temperature in our models.

### 1.4.3 Fourier terms

We used Fourier terms to capture the patterns of seasonality in the data. Fourier terms are a set of sine and cosine functions with different frequencies that can be used to model periodic patterns. For a given period  $P$ , the Fourier terms are defined as follows:

$$\sin_k(t) = \sin\left(\frac{2\pi kt}{P}\right), \quad \cos_k(t) = \cos\left(\frac{2\pi kt}{P}\right)$$

where  $k$  is the frequency and  $t$  is the time. Then a partial Fourier sum can be written as

$$y(t) = \beta_0 + \sum_{k=1}^K \left( \beta_{1k} \sin\left(\frac{2\pi kt}{P}\right) + \beta_{2k} \cos\left(\frac{2\pi kt}{P}\right) \right)$$

where  $K$  is the maximum order of Fourier terms to use.

We used Fourier terms to model the daily and annual seasonality in the data. Since our dataset is missing data for some days, we needed to adjust the  $t$  used to calculate the Fourier terms to account for these gaps. In the `feature_engineering.R` script, we added the variable `counter` to all dataframes. This variable increments by one for each time step (half hour, hour, or day) in the dataframe, reflecting the fact that some days are missing. These `counter` variables were then used later on to generate Fourier terms for our models. We wrote an Rcpp function in our package called `generate_fourier_terms`, which creates a matrix of Fourier terms given a time counter variable representing  $t$ , the maximum order of the Fourier terms  $K$ , and the number of time increments in a period  $P$ .

## 1.5 Scaling

# 2 Simple Regression

## 2.1 Theory

## 2.2 Model

## 2.3 Implementation

# 3 Ridge Regression

## 3.1 Theory

Ridge regression is a method for penalized regression. Consider the model

$$Y_i^0 = \alpha + \beta x_i^0 + \epsilon_i, \quad i = 1, \dots, n$$

where  $\beta \in \mathbb{R}^p$ ,  $\alpha \in \mathbb{R}$ , and for all  $i, l \in \{1, \dots, n\}$ ,  $\mathbb{E}[\epsilon_i] = 0$  and  $\mathbb{E}[\epsilon_i \epsilon_l] = \sigma^2 \delta_{il}$  for some  $\sigma^2 > 0$ . Then the ridge regression estimator is defined as the minimizer of the following objective function:

$$(\hat{\alpha}_\lambda, \hat{\beta}_\lambda) = \operatorname{argmin}_{\alpha \in \mathbb{R}, \beta \in \mathbb{R}^p} \|y^0 - \alpha - X^0 \beta\|_2^2 + \lambda \|\beta\|_2^2$$

where  $\lambda > 0$  is a tuning parameter and  $\|\cdot\|_2$  denotes the Euclidean norm. Ridge regression is thus imposing a penalty on the size of  $\beta$ , with the strength of that penalty determined by the choice of  $\lambda$ . The coefficients will be shrunk towards zero, but will not be set to zero (as opposed to in lasso regression).

## 3.2 Model

## 3.3 Implementation

# 4 Gaussian Process Regression

## 4.1 Theory

A gaussian process  $W = (W(x))_{x \in \mathcal{X}}$  is a collection of random variables, which have a joint Gaussian distribution. One useful fact is that a Gaussian process is completely specified by its mean ( $\mu : \mathcal{X} \rightarrow \mathbb{R}$ ) and covariance ( $k : \mathcal{X}^2 \rightarrow \mathbb{R}$ ) functions - meaning the objective of GPR is to obtain the mean and covariance. This allows us to express any finite collection of  $W$  as follows:

$$(W(x_1), W(x_2), \dots, W(x_n)) \sim N((\mu(x_1), \dots, \mu(x_n)), (k(x_i, x_j))_{ij})$$

We build the following model:

Let  $y_i = f(x_i) + \varepsilon_i$ , where  $\varepsilon_i \sim N(0, \sigma^2)$  and we let  $f$  be a gaussian process. It can be shown that

$$f \mid y^0 \sim \text{GP}(f_n, k_n)$$

where

$$f_n(x) = k^n(x)^T (K + \sigma^2 I_n)^{-1} y^0,$$

$$k(x, x^*) = k(x, x^*) - k^n(x)^T (K + \sigma^2 I_n)^{-1} k^n(x^*)$$

and

$$K = (k(x_i^0, x_j^0))_{ij}$$

$$k^n(x) = (k(x_1^0, x), k(x_n^0, x)).$$

Although if it looks scary just remember a computer can do it! In practice, to find the posterior distribution, we maximise the marginal log-likelihood.

## 4.2 Model

## 4.3 Implementation

For the gaussian process regression, we will use the `kernlab::gausspr` function, which can perform multi-variate gaussian process regression, allowing us to add our additional features as covariates. By convention we will use the rbf kernel:

$$k(x, x') = \exp\left(-\frac{\|x - x'\|^2}{2l}\right)$$

where  $l$  is another hyperparameter to be tuned. Below is our code which fits the gaussian process regression model to given data:

```
gaussian_process_reg <- function(data,
                                class = "DE",
                                kernel = "rbfdot",
                                plot = FALSE,
                                sigma = 100) {

  # Training and test set, first 90% of data is training, last 10% is test
  train_index <- round(nrow(data) * 0.9)
  train_set <- data[1:train_index, ]
  test_set <- data[(train_index + 1):nrow(data), ]

  # Define the Gaussian process model
  gpr_model <- kernlab::gausspr(as.matrix(train_set[, c("toy", "temp")]),
                                as.vector(train_set[[class]]),
                                kernel = kernel, kpar = list(sigma = sigma))

  # Predict the mean function for plotting
  mean_func <- predict(gpr_model, as.vector(data$toy))
  prediction <- data.frame(
    toy = data$toy,
    mean = mean_func)

  data_with_pred <- left_join(data, prediction, by = "toy")

  # Evaluate performance
  test_mean_func <- predict(gpr_model, as.vector(test_set$toy))
  performance <- postResample(pred = test_mean_func, obs = test_set$DE)
```

```

# Include Plots ?
if (plot){
  p1 <- ggplot(data_with_pred, aes(x = toy)) +
    geom_point(aes(y = get(class)), color = "#414141") +
    geom_line(aes(y = mean), color = class_colours[class]) +
    labs(title = "GPR model predictions",
         x = "Date", y = "Demand")
  p1
} else {
  p1 <- NULL
}
return(list(model = gpr_model,
            data <- data_with_pred,
            performance,
            plot = p1))
}

```

Since we are working as bayesians to find the ideal hyperparameters  $l$  and  $\sigma$ , we will minimise the negative log likelihood of the model. On one hand this skirts the issues of cross validation, but on the other hand we are left with a non convex optimisation problem. To get around this issue we outsource to C++, making use of the C++ optimising library `optim`. Then paralling with `RcppArmadillo` to parallelise the optimisation over a two dimensional grid of initial choices. Below is the `RcppArmadillo` code we used to optimise the hyperparameters:

```

#include <RcppArmadillo.h>
#include <optim.hpp>

// [[Rcpp::depends(RcppArmadillo)]]

// Define the RBF kernel
arma::mat rbf_kernel(const arma::mat& X, double l, double sigma) {
  arma::mat K = arma::zeros(X.n_rows, X.n_rows);
  for (unsigned i = 0; i < X.n_rows; ++i) {
    for (unsigned j = 0; j < X.n_rows; ++j) {
      double dist = arma::norm(X.row(i) - X.row(j), 2);
      K(i, j) = std::pow(sigma, 2) * std::exp(-std::pow(dist, 2) / (2 * std::pow(l, 2)));
    }
  }
  return K;
}

// Define the negative log marginal likelihood
double neg_log_marginal_likelihood(const arma::vec& theta, arma::vec* grad_out, void* opt_data) {
  // Extract data and parameters
  arma::mat X = *static_cast<arma::mat*>(opt_data);
  arma::vec y = *static_cast<arma::vec*>(opt_data);
  double l = theta(0);
  double sigma = theta(1);

  // Calculate the kernel matrix
  arma::mat K = rbf_kernel(X, l, sigma);

  // Calculate the log marginal likelihood

```

```

    double log_likelihood = -0.5 * arma::as_scalar(y.t() * arma::solve(K, y)) - 0.5 * arma::log_det(K +

    // Return the negative log marginal likelihood
    return -log_likelihood;
}

// [[Rcpp::export]]
arma::vec optimise_gaussian_process(arma::mat X, arma::vec y) {
    // Initial guess for the parameters
    arma::vec theta = arma::ones(2);

    // Optimise the negative log marginal likelihood
    bool success = optim::de(theta, neg_log_marginal_likelihood, &X);

    // Return the optimised parameters
    return theta;
}

```

## 5 Results